

Defining Relations

DDL statements

Creating Relations in SQL

- Creates the Students relation. Observe that the type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified (domain constraints)
- As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Students  
(sid: CHAR(20),  
name: CHAR(20),  
login: CHAR(10),  
age: INTEGER,  
gpa: REAL)
```

```
CREATE TABLE Enrolled  
(sid: CHAR(20),  
cid: CHAR(20),  
grade: CHAR(2))
```

Destroying and Altering Relations

DROP TABLE Students

- Destroys the relation Students. The schema information *and* the tuples are deleted.

ALTER TABLE Students

ADD COLUMN firstYear: integer

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.

Adding and Deleting Tuples

Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE
FROM Students S
WHERE S.name = 'Smith'
```

Integrity Constraints (ICs)

- IC: condition that must be true for any instance of the database; e.g., domain constraints
 - ICs are specified when schema is defined
 - ICs are checked when relations are modified
- A legal instance of a relation is one that satisfies all specified ICs
 - DBMS should not allow illegal instances
- If the DBMS checks ICs, stored data is more faithful to real-world meaning
 - Avoids data entry errors, too!

Primary Key Constraints

- A set of fields is a key for a relation if :
 1. No two distinct tuples can have same values in all key fields, and
 2. This is not true for any subset of the key.
 - Part 2 false? A *superkey*.
 - If there's >1 key for a relation, one of the keys is chosen (by DBA) to be the *primary key*.
- E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

Primary and Candidate Keys in SQL

- Possibly many candidate keys (specified using UNIQUE), one of which is chosen as the *primary key*.
- ❖ “For a given student and course, there is a single grade.” vs. “Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”
- ❖ Used carelessly, an IC can prevent the storage of database instances that arise in practice!

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid,cid) )
```

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid),  
  UNIQUE (cid, grade) )
```

Foreign Keys, Referential Integrity

- Foreign key : Set of fields in one relation that is used to `refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a `logical pointer'.
- E.g. *sid* is a foreign key referring to **Students**:
 - Enrolled(*sid*: string, *cid*: string, *grade*: string)
 - If all foreign key constraints are enforced, referential integrity is achieved

Foreign Keys in SQL

- Only students listed in the Students relation should be allowed to enroll for courses

```
CREATE TABLE Enrolled
(sid CHAR(20), cid CHAR(20), grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Enforcing Referential Integrity

- Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students
- What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a Students tuple is deleted?
 - Also delete all Enrolled tuples that refer to it
 - Disallow deletion of a Students tuple that is referred to
 - Set *sid* in Enrolled tuples that refer to it to a *default sid*
 - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting '*unknown*' or '*inapplicable*') but *sid* is part of a primary key.
 - Default is NO ACTION ie DELETE or UPDATE is rejected
- **Similar if primary key value of Students tuple is updated**

Referential Integrity in SQL

- SQL/92 and SQL:1999 support all 4 options on deletes and updates
 - Default is **NO ACTION**
(*delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to deleted tuple)
 - **SET NULL / SET DEFAULT**
(sets foreign key value of referencing tuple)
 - sid (CHAR (20) DEFAULT '53995')


```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

Where do ICs Come From?

- ICs are based upon the semantics of the real-world enterprise that is being described in the database relations
- We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance
 - An IC is a statement about *all possible* instances!
- Key and foreign key ICs are the most common; more general ICs supported too



Logical DB Design: ER to Relational

- Entity sets to tables:



```
CREATE TABLE Employees  
  (ssn CHAR(11),  
   name CHAR(20),  
   lot INTEGER,  
   PRIMARY KEY (ssn))
```

ER Model Basics (2)




- *Relationship*: Association among two or more entities. E.g., Rajendra works in research & development department.
- *Relationship Set*: Collection of similar relationships.
 - An n-ary relationship set R relates n entity sets $E_1 \dots E_n$; each relationship in R involves entities $e_1 \in E_1, \dots, e_n \in E_n$
 - Same entity set could participate in different relationship sets, or in different “roles” in same set.

Relationship Sets to Tables


- In translating a relationship set to a relation/table, attributes of the relation must include:
 - Keys for each participating entity set (as foreign keys)
 - This set of attributes forms a superkey for the relation
 - All descriptive attributes

```
CREATE TABLE Works_In(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees,  
    FOREIGN KEY (did)  
        REFERENCES Departments)
```


Review: Key Constraints




- Each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many



Many-to-1



Many-to-Many

Translation to relational model?

Department-Manages-Employee

```
CREATE TABLE Departments
    (dname CHAR(20),
     did INTEGER,
     budget INTEGER,
     PRIMARY KEY (did))
```

```
CREATE TABLE Employees
    (ssn CHAR(11),
     name CHAR(20),
     lot INTEGER,
     PRIMARY KEY (ssn))
```

Translating ER Diagrams with Key Constraints


- Map relationship to a table:
 - Note that `did` is the key now!
 - Separate tables for Employees and Departments
- Since each department has a unique manager, we could instead combine Manages and Departments

```
CREATE TABLE Manages(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    FOREIGN KEY (did) REFERENCES Departments)
```

```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    name CHAR(20),  
    budget REAL,  
    ssn CHAR(11),  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees)
```

Review: Participation Constraints

- Does every department have a manager?
 - If so, this is a *participation constraint*: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)
 - Participation: every dept has manager
 - Key: at most one manager



Participation Constraints in SQL

- We can capture participation constraints involving one entity set in a binary relationship


```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE NO ACTION)
```

Contd...

- On delete no actionis default and need not be explicitly specified
- Ensures that an Employee tuple can not be deleted while it is pointed to by a Dept-Mgr tuple
- If we wish to delete that Employee tuple we first need to have a new employee as manager in Dept-Mgr
- If we specify CASCADE instead of NO ACTION
 - Deleting information about the department itself because the manager has been fired !!!!
- The constraint that every dept has a manager cant be captured using first option (separate manages relation option)
- If NOTNULL is added to ssn and did
 - It would prevent firing a manager

Review: Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.



Translating Weak Entity Sets

- Weak entity set and identifying relationship set (course) are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.


```
CREATE TABLE Dep_Policy (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

Reports-To Relationship Set

```
CREATE TABLE Reports-To(  
  supervisor_ssn CHAR(11),  
  subordinate_ssn CHAR(11),  
  PRIMARY KEY (supervisor_ssn,subordinate_ssn),  
  FOREIGN KEY (supervisor_ssn) REFERENCES Employees (ssn),  
  FOREIGN KEY (subordinate_ssn) REFERENCES Employees (ssn))
```


Review: ISA Hierarchies

- ❖ As in C++, or other PLs, attributes are inherited.
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.




- *Overlap constraints*: Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*)
- *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*)


Translating ISA Hierarchies to Relations

- **General approach:**
 - 3 relations: Employees, Hourly_Emps and Contract_Emps.
 - *Hourly_Emps*:
 - Every employee is recorded in Employees For hourly emps, extra info recorded in Hourly_Emps (*hourly_wages*, *hours_worked*, *ssn*)
 - must delete Hourly_Emps tuple if referenced Employees tuple is deleted)
 - Queries involving all employees easy, those involving just Hourly_Emps require a join to get some attributes.
- **Alternative: Just Hourly_Emps and Contract_Emps.**
 - *Hourly_Emps*: *ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*.
 - Each employee must be in one of these two subclasses.

Review: Binary vs. Ternary Relationships



Bad design



Better design

Translating to Relational Model

- Deletion of Employee leads to deletion of all policies owned by employee and all dependents who are beneficiaries of those policies
- Each dependent is required to have a covering policy (implicit NOT NULL)

Binary vs. Ternary Relationships (Contd.)


- The key constraints allow us to combine Purchaser with Policies and Beneficiary with Dependents.
- Participation constraints lead to **NOT NULL** constraints.
- What if Policies is a weak entity set?

```
CREATE TABLE Policies (  
  policyid INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (policyid).  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (  
  pname CHAR(20),  
  age INTEGER,  
  policyid INTEGER,  
  PRIMARY KEY (pname, policyid).  
  FOREIGN KEY (policyid) REFERENCES Policies,  
  ON DELETE CASCADE)
```

Aggregation

- Used when we have to model a relationship involving (entity sets and) a relationship set
 - **Aggregation** allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships



☛ *Aggregation vs. ternary relationship:*

- ❖ Monitors is a distinct relationship, with a descriptive attribute
- ❖ Also, can say that each sponsorship is monitored by at most one employee

Aggregation

- Projects, Departments, Sponsors translated as discussed
 - Projects(pid, started_on, pbudget)
 - Departments (did, dname, budget)
 - Sponsors (pid,did,since)
- Monitors (until, pid, did, ssn)
- Employee (ssn, name, lot)
- Why Sponsors is required?
 - To record descriptive attributes of Sponsor (since)
 - Not every sponsorship has a monitor and thus some (pid,did) pairs in sponsors may not appear in Monitors
- If Sponsors
 - has no descriptive attributes
 - Has total participation in Monitors
- Then every possible instance of Sponsors can be obtained from (pid,did) columns of Monitors (Sponsors can be dropped)

Views

- A view is just a relation, but we store a *definition*, rather than a set of tuples]
- From students and Enrolled relations

```
CREATE VIEW YoungActiveStudents (name, grade)
AS SELECT S.name, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age<21
```

- ❖ Views can be dropped using the **DROP VIEW** command
 - How to handle **DROP TABLE** if there's a view on the table?
 - DROP TABLE command has options to let the user specify this

Updating Views

- Updates to be specified only on views that are defined on a single base table (SQL 92) using just projection and selection, with no use of aggregate operations (updatable views)

```
CREATE TABLE Good Students (sid,gpa)  
AS SELECT S.sid, S.gpa  
FROM Students S  
WHERE S.gpa > 9.0
```

If view does not contain key of the base table, several rows in table would correspond to a single row in view

- SQL 99 allows updates on views using more than one table where the PK of that table is included in the fields of view
- Views defined using UNION, INTERSECT, EXCEPT cant be inserted into even if they are updatable

Insertion

- If we insert a GoodStudents row by inserting a row into Students, using null values in columns of students that do not appear in GoodStudents(eg sams,login)
- Note that PK columns are not allowed to contain null values
- Therefore if we attempt to insert rows through a view that does not contain the PK of the underlying table, the insertions will be rejected

Views and Security

- Views can be used to present necessary information (or a summary), while hiding details in underlying relation(s)
- Given YoungStudents, but not Students or Enrolled, we can find students *s* who are enrolled, but not the *cid*'s of the courses they are enrolled in

Relational Model: Summary

- A tabular representation of data
- Simple and intuitive, currently the most widely used
- Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations
 - Two important ICs: primary and foreign keys
 - In addition, we *always* have domain constraints
- Powerful and natural query languages exist
- Rules to translate ER to relational model